



Atomic compare-and-Swap (CAS) instructions (Zacas)

Version 0.1, 4/2023: This document is in development state. See <http://riscv.org/spec-state> for details.

Table of Contents

Preamble.....	1
Copyright and license information.....	2
Contributors.....	3
1. Introduction.....	4
2. Word/Doubleword/Quadword CAS (AMOCAS.W/D/Q).....	5
3. Additional AMO PMAs.....	7

Preamble



This document is in the [Development state](#)

Assume everything can change. This draft specification will change before being accepted as standard, so implementations made to this draft specification will likely not conform to the future standard.

Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Copyright 2022 by RISC-V International.

Contributors

This RISC-V specification has been contributed to directly or indirectly by: Greg Favor, Ved Shanbhogue


Chapter 1. Introduction

Compare-and-swap (CAS) provides an easy and typically faster way to perform thread synchronization operations when supported as a hardware instruction. CAS is typically used by lock-free and wait-free algorithms. This extension proposes CAS instructions to operate on 32-bit, 64-bit, and 128-bit (RV64 only) data values. The CAS instruction supports the C++11 atomic compare and exchange operation.

While compare-and-swap for XLEN wide data may be accomplished using LR/SC, the CAS atomic instructions scale better to highly parallel systems than LR/SC. Many lock-free algorithms, such as a lock-free queue, require manipulation of pointer variables. A simple CAS operation may not be sufficient to guard against what is commonly referred to as the ABA problem in such algorithms that manipulate pointer variables. To avoid the ABA problem, the algorithms associate a reference counter with the pointer variable and perform updates using a quadword compare and swap (of both the pointer and the counter). The double and quadword CAS instructions support implementation of algorithms for ABA problem avoidance.

The Zacas extension depends upon the A extension.

Chapter 2. Word/Doubleword/Quadword CAS (AMOCAS.W/D/Q)



AMOCAS.W atomically loads 32-bits of a data value from address in **rs1**, compares the loaded value to a 32-bit value held in **rd** and if the comparison is bitwise equal, then stores the 32-bit value held in **rs2** to the original address in **rs1**. The value loaded from memory is placed into register **rd**. For RV64, **AMOCAS.W** always sign-extends the value placed in **rd**, and ignores the upper 32 bits of the original value in **rd** and **rs2**. The operation performed by **AMOCAS.W** is as follows:

```
temp = *[rs1]
if temp == [rd]
    *[rs1] = [rs2]
endif
[rd] = temp
```

AMOCAS.D is similar to **AMOCAS.W** but operates on 64-bit data values.

For RV32, **AMOCAS.D** atomically loads 64-bits of a data value from address in **rs1**, compares the loaded value to a 64-bits value held in a register pair consisting of **rd** and **rd+1** and if the comparison is bitwise equal, then stores the 64-bit value held in the register pair **rs2** and **rs2+1** to the original address in **rs1**. The value loaded from memory is placed into the register pair **rd** and **rd+1**. The instruction requires the first register in the pair to be even numbered; encodings with odd numbered registers specified in **rs2** and **rd** are reserved. When the first register of a source register pair is **x0**, then both halves of the pair read as zero. When the first register of a destination register pair is **x0**, then the entire register result is discarded and neither destination register is written. The operation performed by **AMOCAS.D** for RV32 is as follows:

```
temp0 = *([rs1]+0)
temp1 = *([rs1]+4)
comp0 = (rd == x0) ? 0 : [rd];
comp1 = (rd == x0) ? 0 : [rd+1];
swap0 = (rs2 == x0) ? 0 : [rs2];
swap1 = (rs2 == x0) ? 0 : [rs2+1];
If (temp0 == comp0) && (temp1 == comp1)
    *([rs1]+0) = swap0
    *([rs1]+4) = swap1
endif
if ( rd != x0 )
    [rd]    = temp0
    [rd+1] = temp1
endif
```

For RV64, **AMOCAS.D** atomically loads 64-bits of a data value from address in **rs1**, compares the loaded value to a 64-bit value held in **rd** and if the comparison is bitwise equal, then stores the 64-bit value held in **rs2** to the original address in **rs1**. The value loaded from memory is placed into register **rd**. The operation performed by **AMOCAS.D** for RV64 is as follows:

```
temp = *[rs1]
if temp == [rd]
    *[rs1] = [rs2]
endif
[rd] = temp
```

AMOCAS.Q (RV64 only) atomically loads 128-bits of a data value from address in **rs1**, compares the loaded value to a 128-bits value held in a register pair consisting of **rd** and **rd+1** and if the comparison is bitwise equal, then stores the 128-bit value held in the register pair **rs2** and **rs2+1** to the original address in **rs1**. The value loaded from memory is placed into the register pair **rd** and **rd+1**. The instruction requires the first register in the pair to be even numbered; encodings with odd numbered registers specified in **rs2** and **rd** are reserved. When the first register of a source register pair is **x0**, then both halves of the pair read as zero. When the first register of a destination register pair is **x0**, then the entire register result is discarded and neither destination register is written. The operation performed by **AMOCAS.Q** is as follows:

```
temp0 = *([rs1]+0)
temp1 = *([rs1]+8)
comp0 = (rd == x0) ? 0 : [rd];
comp1 = (rd == x0) ? 0 : [rd+1];
swap0 = (rs2 == x0) ? 0 : [rs2];
swap1 = (rs2 == x0) ? 0 : [rs2+1];
If (temp0 == comp0) && (temp1 == comp1)
    *([rs1]+0) = swap0
    *([rs1]+8) = swap1
endif
if ( rd != x0 )
    [rd] = temp0
    [rd+1] = temp1
endif
```



For a future RV128 extension, **AMOCAS.Q** would encode a single XLEN=128 register in **rs2** and **rd**.

Just as for AMOs in the A extension, **AMOCAS.W/D/Q** requires that the address held in **rs1** be naturally aligned to the size of the operand (i.e., 16-byte aligned for 128-bit words, eight-byte aligned for 64-bit words, and four-byte aligned for 32-bit words). And the same exception options apply if the address is not naturally aligned.

Just as for AMOs in the A extension, the **AMOCAS.W/D/Q** optionally provide release consistency semantics, using the **aq** and **rl** bits, to help implement multiprocessor synchronization.

Chapter 3. Additional AMO PMAs

There are four levels of PMA support defined for AMOs in the A extension. Zacas defines three additional levels of support: `AMOCasW`, `AMOCasD`, and `AMOCasQ`.

`AMOCasW` indicates that in addition to instructions indicated by `AMOArithmetic` level support, the `AMOCAS.W` instruction is supported. `AMOCasD` indicates that in addition to instructions indicated by `AMOCasW` level support, the `AMOCAS.D` instruction is supported. `AMOCasQ` indicates that all RISC-V AMOs are supported.



`AMOCasW/D/Q` require `AMOArithmetic` level support as the `AMOCAS.W/D/Q` instructions require ability to perform an arithmetic comparison and a swap operation.